

In [2]: *#upload tiny imagenet folder into jupyter project*

```
#import zipfile as zf  
#files = zf.ZipFile("tiny-imagenet-200.zip", 'r')  
#files.extractall()  
#files.close
```

In [3]: *%matplotlib inline*

```
import matplotlib.pyplot as plt  
import numpy as np  
import torch  
import torchvision  
import torchvision.datasets as datasets  
import torch.utils.data as data  
from torchvision.utils import make_grid  
import torchvision.transforms as transforms  
import torch.nn as nn  
import torch.nn.functional as F  
import torch.optim as optim  
  
import os  
  
import vgg  
import resnet  
import googlenet  
import alexnet
```

In [4]: *# If there are GPUs, choose the first one for computing. Otherwise use CPU.*

```
device = torch.device("cuda:0" if torch.cuda.is_available() else "cpu")  
print(device)  
# If 'cuda:0' is printed, it means GPU is available.
```

cuda:0

```

In [5]: data_transforms = {
    'train': transforms.Compose([
        transforms.RandomRotation(20),
        transforms.RandomHorizontalFlip(0.5),
        transforms.ToTensor(),
        transforms.Normalize([0.4802, 0.4481, 0.3975], [0.2302, 0.2265, 0.2262])
    ]),
    'val': transforms.Compose([
        transforms.ToTensor(),
        transforms.Normalize([0.4802, 0.4481, 0.3975], [0.2302, 0.2265, 0.2262])
    ]),
    'test': transforms.Compose([
        transforms.ToTensor(),
        transforms.Normalize([0.4802, 0.4481, 0.3975], [0.2302, 0.2265, 0.2262])
    ])
}

data_dir = 'tiny-imagenet-200/'

num_workers = {
    'train' : 100,
    'val'    : 0,
    'test'   : 0
}


image_datasets = {x: datasets.ImageFolder(os.path.join(data_dir, x),
                                                    data_transforms[x])
                  for x in ['train', 'val', 'test']}
dataloaders = {x: torch.utils.data.DataLoader(image_datasets[x], batch_size=100,
                                                    shuffle=True, num_workers=num_workers,
                                                    for x in ['train', 'val', 'test'])
              for x in ['train', 'val', 'test']}
dataset_sizes = {x: len(image_datasets[x]) for x in ['train', 'val', 'test']}

```

```
In [8]: def imshow(img):
        img = img.numpy().transpose((1, 2, 0))
        img = np.clip(img, 0, 1)
        plt.imshow(img)

        images, labels = next(iter(dataloaders['val']))
        print(labels)
        grid = make_grid(images[:4], nrow=4)
        imshow(grid)
```

```
tensor([0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0,
        0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0,
        0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0,
        0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0,
        0, 0, 0, 0])
```



```
In [22]: net = vgg.vgg16()    # Create the network instance.
net.to(device)
```

```
Out[22]: VGG(
  (features): Sequential(
    (0): Conv2d(3, 64, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
    (1): ReLU(inplace=True)
    (2): Conv2d(64, 64, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
    (3): ReLU(inplace=True)
    (4): MaxPool2d(kernel_size=2, stride=2, padding=0, dilation=1, ceil_mode=False)
  )
  (5): Conv2d(64, 128, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
  (6): ReLU(inplace=True)
  (7): Conv2d(128, 128, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
  (8): ReLU(inplace=True)
  (9): MaxPool2d(kernel_size=2, stride=2, padding=0, dilation=1, ceil_mode=False)
  (10): Conv2d(128, 256, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
  (11): ReLU(inplace=True)
  (12): Conv2d(256, 256, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
  (13): ReLU(inplace=True)
  (14): Conv2d(256, 256, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
  (15): ReLU(inplace=True)
  (16): MaxPool2d(kernel_size=2, stride=2, padding=0, dilation=1, ceil_mode=False)
  (17): Conv2d(256, 512, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
  (18): ReLU(inplace=True)
  (19): Conv2d(512, 512, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
  (20): ReLU(inplace=True)
  (21): Conv2d(512, 512, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
  (22): ReLU(inplace=True)
  (23): MaxPool2d(kernel_size=2, stride=2, padding=0, dilation=1, ceil_mode=False)
  (24): Conv2d(512, 512, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
  (25): ReLU(inplace=True)
  (26): Conv2d(512, 512, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
  (27): ReLU(inplace=True)
  (28): Conv2d(512, 512, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
  (29): ReLU(inplace=True)
  (30): MaxPool2d(kernel_size=2, stride=2, padding=0, dilation=1, ceil_mode=False)
  )
  (avgpool): AdaptiveAvgPool2d(output_size=(7, 7))
  (classifier): Sequential(
    (0): Linear(in_features=25088, out_features=4096, bias=True)
    (1): ReLU(inplace=True)
    (2): Dropout(p=0.5, inplace=False)
    (3): Linear(in_features=4096, out_features=4096, bias=True)
    (4): ReLU(inplace=True)
    (5): Dropout(p=0.5, inplace=False)
    (6): Linear(in_features=4096, out_features=1000, bias=True)
  )
)
```

```
In [11]: # We use cross-entropy as loss function.  
loss_func = nn.CrossEntropyLoss()  
# We use stochastic gradient descent (SGD) as optimizer.  
opt = optim.SGD(net.parameters(), lr=0.001, momentum=0.9)
```

```

In [23]: avg_losses = []    # Avg. Losses.
         epochs = 11        # Total epochs.
         print_freq = 100   # Print frequency.

         for epoch in range(epochs): # Loop over the dataset multiple times.
             running_loss = 0.0      # Initialize running loss.
             for i, data in enumerate(dataloaders['train'], 0):

                 net.train()

                 # Get the inputs
                 inputs, labels = data

                 # Move the inputs to the specified device.
                 inputs, labels = inputs.to(device), labels.to(device)

                 # Zero the parameter gradients.
                 opt.zero_grad()

                 # Forward step.
                 outputs = net(inputs)
                 loss = loss_func(outputs, labels)

                 # Backward step.
                 loss.backward()

                 # Optimization step (update the parameters).
                 opt.step()

                 # Print statistics.
                 running_loss += loss.item()
                 if i % print_freq == print_freq - 1: # Print every several mini-batches.
                     avg_loss = running_loss / print_freq
                     print('[epoch: {}, i: {:5d}] avg mini-batch loss: {:.3f}'.format(
                         epoch, i, avg_loss))
                     avg_losses.append(avg_loss)
                     running_loss = 0.0

             print('Finished Training.')

```

```

[epoch: 0, i:    99] avg mini-batch loss: 6.924
[epoch: 0, i:   199] avg mini-batch loss: 6.924
[epoch: 0, i:   299] avg mini-batch loss: 6.926
[epoch: 0, i:   399] avg mini-batch loss: 6.925
[epoch: 0, i:   499] avg mini-batch loss: 6.925
[epoch: 0, i:   599] avg mini-batch loss: 6.926
[epoch: 0, i:   699] avg mini-batch loss: 6.924
[epoch: 0, i:   799] avg mini-batch loss: 6.925
[epoch: 0, i:   899] avg mini-batch loss: 6.925
[epoch: 0, i:   999] avg mini-batch loss: 6.925
[epoch: 1, i:    99] avg mini-batch loss: 6.923
[epoch: 1, i:   199] avg mini-batch loss: 6.927
[epoch: 1, i:   299] avg mini-batch loss: 6.924
[epoch: 1, i:   399] avg mini-batch loss: 6.923
[epoch: 1, i:   499] avg mini-batch loss: 6.922
[epoch: 1, i:   599] avg mini-batch loss: 6.924

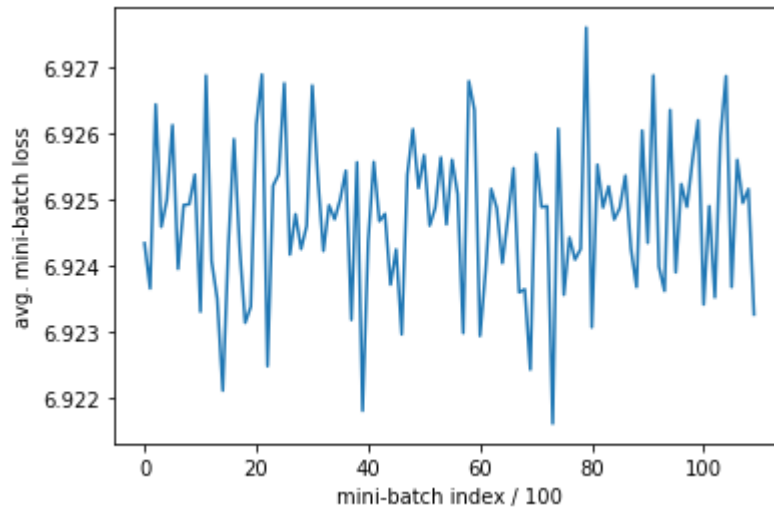
```

```
[epoch: 1, i: 699] avg mini-batch loss: 6.926
[epoch: 1, i: 799] avg mini-batch loss: 6.924
[epoch: 1, i: 899] avg mini-batch loss: 6.923
[epoch: 1, i: 999] avg mini-batch loss: 6.923
[epoch: 2, i: 99] avg mini-batch loss: 6.926
[epoch: 2, i: 199] avg mini-batch loss: 6.927
[epoch: 2, i: 299] avg mini-batch loss: 6.922
[epoch: 2, i: 399] avg mini-batch loss: 6.925
[epoch: 2, i: 499] avg mini-batch loss: 6.925
[epoch: 2, i: 599] avg mini-batch loss: 6.927
[epoch: 2, i: 699] avg mini-batch loss: 6.924
[epoch: 2, i: 799] avg mini-batch loss: 6.925
[epoch: 2, i: 899] avg mini-batch loss: 6.924
[epoch: 2, i: 999] avg mini-batch loss: 6.925
[epoch: 3, i: 99] avg mini-batch loss: 6.927
[epoch: 3, i: 199] avg mini-batch loss: 6.925
[epoch: 3, i: 299] avg mini-batch loss: 6.924
[epoch: 3, i: 399] avg mini-batch loss: 6.925
[epoch: 3, i: 499] avg mini-batch loss: 6.925
[epoch: 3, i: 599] avg mini-batch loss: 6.925
[epoch: 3, i: 699] avg mini-batch loss: 6.925
[epoch: 3, i: 799] avg mini-batch loss: 6.923
[epoch: 3, i: 899] avg mini-batch loss: 6.926
[epoch: 3, i: 999] avg mini-batch loss: 6.922
[epoch: 4, i: 99] avg mini-batch loss: 6.924
[epoch: 4, i: 199] avg mini-batch loss: 6.926
[epoch: 4, i: 299] avg mini-batch loss: 6.925
[epoch: 4, i: 399] avg mini-batch loss: 6.925
[epoch: 4, i: 499] avg mini-batch loss: 6.924
[epoch: 4, i: 599] avg mini-batch loss: 6.924
[epoch: 4, i: 699] avg mini-batch loss: 6.923
[epoch: 4, i: 799] avg mini-batch loss: 6.925
[epoch: 4, i: 899] avg mini-batch loss: 6.926
[epoch: 4, i: 999] avg mini-batch loss: 6.925
[epoch: 5, i: 99] avg mini-batch loss: 6.926
[epoch: 5, i: 199] avg mini-batch loss: 6.925
[epoch: 5, i: 299] avg mini-batch loss: 6.925
[epoch: 5, i: 399] avg mini-batch loss: 6.926
[epoch: 5, i: 499] avg mini-batch loss: 6.925
[epoch: 5, i: 599] avg mini-batch loss: 6.926
[epoch: 5, i: 699] avg mini-batch loss: 6.925
[epoch: 5, i: 799] avg mini-batch loss: 6.923
[epoch: 5, i: 899] avg mini-batch loss: 6.927
[epoch: 5, i: 999] avg mini-batch loss: 6.926
[epoch: 6, i: 99] avg mini-batch loss: 6.923
[epoch: 6, i: 199] avg mini-batch loss: 6.924
[epoch: 6, i: 299] avg mini-batch loss: 6.925
[epoch: 6, i: 399] avg mini-batch loss: 6.925
[epoch: 6, i: 499] avg mini-batch loss: 6.924
[epoch: 6, i: 599] avg mini-batch loss: 6.925
[epoch: 6, i: 699] avg mini-batch loss: 6.925
[epoch: 6, i: 799] avg mini-batch loss: 6.924
[epoch: 6, i: 899] avg mini-batch loss: 6.924
[epoch: 6, i: 999] avg mini-batch loss: 6.922
[epoch: 7, i: 99] avg mini-batch loss: 6.926
[epoch: 7, i: 199] avg mini-batch loss: 6.925
[epoch: 7, i: 299] avg mini-batch loss: 6.925
```

```
[epoch: 7, i: 399] avg mini-batch loss: 6.922
[epoch: 7, i: 499] avg mini-batch loss: 6.926
[epoch: 7, i: 599] avg mini-batch loss: 6.924
[epoch: 7, i: 699] avg mini-batch loss: 6.924
[epoch: 7, i: 799] avg mini-batch loss: 6.924
[epoch: 7, i: 899] avg mini-batch loss: 6.924
[epoch: 7, i: 999] avg mini-batch loss: 6.928
[epoch: 8, i: 99] avg mini-batch loss: 6.923
[epoch: 8, i: 199] avg mini-batch loss: 6.926
[epoch: 8, i: 299] avg mini-batch loss: 6.925
[epoch: 8, i: 399] avg mini-batch loss: 6.925
[epoch: 8, i: 499] avg mini-batch loss: 6.925
[epoch: 8, i: 599] avg mini-batch loss: 6.925
[epoch: 8, i: 699] avg mini-batch loss: 6.925
[epoch: 8, i: 799] avg mini-batch loss: 6.924
[epoch: 8, i: 899] avg mini-batch loss: 6.924
[epoch: 8, i: 999] avg mini-batch loss: 6.926
[epoch: 9, i: 99] avg mini-batch loss: 6.924
[epoch: 9, i: 199] avg mini-batch loss: 6.927
[epoch: 9, i: 299] avg mini-batch loss: 6.924
[epoch: 9, i: 399] avg mini-batch loss: 6.924
[epoch: 9, i: 499] avg mini-batch loss: 6.926
[epoch: 9, i: 599] avg mini-batch loss: 6.924
[epoch: 9, i: 699] avg mini-batch loss: 6.925
[epoch: 9, i: 799] avg mini-batch loss: 6.925
[epoch: 9, i: 899] avg mini-batch loss: 6.926
[epoch: 9, i: 999] avg mini-batch loss: 6.926
[epoch: 10, i: 99] avg mini-batch loss: 6.923
[epoch: 10, i: 199] avg mini-batch loss: 6.925
[epoch: 10, i: 299] avg mini-batch loss: 6.924
[epoch: 10, i: 399] avg mini-batch loss: 6.926
[epoch: 10, i: 499] avg mini-batch loss: 6.927
[epoch: 10, i: 599] avg mini-batch loss: 6.924
[epoch: 10, i: 699] avg mini-batch loss: 6.926
[epoch: 10, i: 799] avg mini-batch loss: 6.925
[epoch: 10, i: 899] avg mini-batch loss: 6.925
[epoch: 10, i: 999] avg mini-batch loss: 6.923
Finished Training.
```



```
In [24]: plt.plot(avg_losses)
plt.xlabel('mini-batch index / {}'.format(print_freq))
plt.ylabel('avg. mini-batch loss')
plt.show()
```



```
In [25]: # Get test accuracy.
correct = 0
total = 0
with torch.no_grad():
    for i, data in enumerate(dataloaders['test']):
        net.eval()
        images, labels = data
        images, labels = images.to(device), labels.to(device)
        outputs = net(images)
        _, predicted = torch.max(outputs.data, 1)
        total += labels.size(0)
        correct += (predicted == labels).sum().item()

print('Accuracy of the network on the 10000 test images: %d %%' % (
    100 * correct / total))
```

Accuracy of the network on the 10000 test images: 23 %

In []: